


Technical Report: End-to-End Customer Churn Data Engineering Pipeline

Executive Summary and Pipeline Overview

This project implements an end-to-end **Customer Churn Analytics and Prediction pipeline**. Customer data is sourced from **MariaDB** and supporting files (`customer_support_tickets.csv`, `offers.csv`, and `usage.json`) and ingested into **HDFS** using **Apache NiFi**.

The raw data is then processed and transformed using **PySpark**, with the curated datasets stored back in HDFS. **Apache Hive** is subsequently used to build the Data Warehouse, providing structured data for **reporting and business analysis**. The prepared data also serves as the foundation for developing a **Machine Learning model for customer churn prediction**.



1. Data Ingestion Pipeline — Apache NiFi to HDFS

Apache NiFi is used as the data ingestion layer to collect data from both database and file-based sources and transfer it into HDFS. Two ingestion patterns are implemented:

- **database ingestion** for the customer data
- **file ingestion** for the remaining datasets.

1. MariaDB Customer Data

The customer data is stored in a **MariaDB database**. NiFi connects to MariaDB using the QueryDatabaseTable processor together with a configured **DBCPCConnectionPool Controller Service**.

The connection is established by configuring the MariaDB JDBC driver, database URL, username, and password within the controller service. The QueryDatabaseTable processor then uses this connection to query the customer table and generate FlowFiles containing the retrieved records.

The resulting data passes through ConvertRecord to ensure the required record format before being written to HDFS using PutHDFS.

The screenshot shows the 'Controller Service Details' window for the 'DBCPCConnectionPool' service. The 'PROPERTIES' tab is active, displaying a table of configuration properties. The 'Database Connection URL' is set to 'jdbc:mysql://localhost:3306/churn_src?useSSL=...'. The 'Database Driver Class Name' is 'org.mariadb.jdbc.Driver'. The 'Database Driver Location(s)' is '/usr/local/nifi/nifi-1.14.0/lib/mariadb-java-client-...'. The 'Kerberos Credentials Service' is 'No value set'. The 'Kerberos Principal' is 'No value set'. The 'Kerberos Password' is 'No value set'. The 'Database User' is 'root'. The 'Password' is 'Sensitive value set'. The 'Max Wait Time' is '500 millis'. The 'Max Total Connections' is '8'. The 'Validation query' is 'No value set'. The 'Minimum Idle Connections' is '0'. The 'Max Idle Connections' is '8'. The 'Max Connection Lifetime' is '-1'. An 'OK' button is at the bottom right.

Property	Value
Database Connection URL	jdbc:mysql://localhost:3306/churn_src?useSSL=...
Database Driver Class Name	org.mariadb.jdbc.Driver
Database Driver Location(s)	/usr/local/nifi/nifi-1.14.0/lib/mariadb-java-client-...
Kerberos Credentials Service	No value set
Kerberos Principal	No value set
Kerberos Password	No value set
Database User	root
Password	Sensitive value set
Max Wait Time	500 millis
Max Total Connections	8
Validation query	No value set
Minimum Idle Connections	0
Max Idle Connections	8
Max Connection Lifetime	-1

- **Connection Target:** Connects NiFi to the **churn_src** MariaDB database on **localhost:3306**.
- **Driver & Library:** Uses the **org.mariadb.jdbc.Driver** class with the driver JAR located at **/usr/local/nifi/nifi-1.14.0/lib/mariadb-java-client-...**
- **Authentication:** Authenticates using the **root** user account.
- **Pool Limits:** Caps total active and idle connections at **8** to manage database load.
- **Timeout:** Limits connection wait time to **500 ms** before timing out.

QueryDatabaseTable processor

The screenshot shows the 'Configure Processor' window for the 'QueryDatabaseTable' processor. The 'PROPERTIES' tab is active, displaying a table of configuration properties. The 'Database Connection Pooling Service' is 'DBCPCConnectionPool'. The 'Database Type' is 'MySQL'. The 'Table Name' is 'stg_customers'. The 'Columns to Return' is 'No value set'. The 'Additional WHERE clause' is 'No value set'. The 'Custom Query' is 'No value set'. The 'Maximum-value Columns' is 'customer_id'. The 'Max Wait Time' is '0 seconds'. The 'Fetch Size' is '1000'. The 'Max Rows Per Flow File' is '1000'. The 'Output Batch Size' is '1'. The 'Maximum Number of Fragments' is '0'. 'CANCEL' and 'APPLY' buttons are at the bottom.

Property	Value
Database Connection Pooling Service	DBCPCConnectionPool
Database Type	MySQL
Table Name	stg_customers
Columns to Return	No value set
Additional WHERE clause	No value set
Custom Query	No value set
Maximum-value Columns	customer_id
Max Wait Time	0 seconds
Fetch Size	1000
Max Rows Per Flow File	1000
Output Batch Size	1
Maximum Number of Fragments	0

- **Database Connection Pooling Service:** DBCPCConnectionPool ;Links the processor to the database connection pool defined earlier.
- **Database Type:** MySQL ;Specifies the SQL dialect used for query generation.
- **Table Name:** stg_customers ;Target source table to extract records from.
- **Maximum-value Columns:** customer_id Tracks the highest customer_id state across executions to ensure only new or updated records are incrementally ingested.
- **Fetch Size:** 1000 ;Retrieves 1,000 rows per JDBC network call to balance memory usage and extraction speed.
- **Max Rows Per Flow File:** 1000 ;Splits extracted data into chunks of 1,000 records per output FlowFile.
- **Output Batch Size:** 1 ;Emits 1 FlowFile per processor execution cycle.

ConvertRecord Processor

Configure Processor

Stopped

SETTINGS SCHEDULING PROPERTIES COMMENTS

Required field

Property	Value
Record Reader	AvroReader
Record Writer	CSVRecordSetWriter
Include Zero Record FlowFiles	false

CANCEL APPLY

- **Purpose:** Converts binary Avro data from QueryDatabaseTable into plain CSV format before writing to HDFS.
- **Record Reader (AvroReader):** Reads incoming FlowFiles using the embedded Avro schema.
- **Record Writer (CSVRecordSetWriter):** Formats the parsed records into comma-separated text rows.
- **Include Zero Record FlowFiles (false):** Prevents creation of empty CSV files in HDFS when queries return 0 rows.

PutHDFS Processor

Configure Processor

Stopped

SETTINGS SCHEDULING PROPERTIES COMMENTS

Required field

Property	Value
Hadoop Configuration Resources	/home/hadoop/hadoop/etc/hadoop/core-site.xml,/home/...
Kerberos Credentials Service	No value set
Kerberos Principal	No value set
Kerberos Keytab	No value set
Kerberos Password	No value set
Kerberos Relogin Period	4 hours
Additional Classpath Resources	No value set
Directory	/churn_proj/customers
Conflict Resolution Strategy	replace
Block Size	No value set
IO Buffer Size	No value set
Replication	No value set

CANCEL APPLY

- **Purpose:** Writes output CSV files directly to the HDFS storage cluster.
- **Hadoop Configuration Resources:** Points to Hadoop configuration files (core-site.xml, hdfs-site.xml) to establish cluster connection settings.
- **Directory (/churn_proj/customers):** Specifies the destination path in HDFS where customer CSV files will be stored.
- **Conflict Resolution Strategy (replace):** Overwrites existing files of the same name in HDFS to maintain updated table data.

Pipeline flow:

MariaDB (customer table) → QueryDatabaseTable → ConvertRecord → PutHDFS → HDFS



```

[hue@localhost ~]$ hdfs dfs -ls /churn_proj/customers
Found 11 items
-rw-r--r-- 1 root hadoop 111307 2026-09-02 08:36 /churn_proj/customers/0f3
eee39-977c-421b-9cf9-c85479f78b8d
-rw-r--r-- 1 root hadoop 16962 2026-09-02 08:36 /churn_proj/customers/1dc
511f3-dd74-4f7c-85ac-86c405772f15
-rw-r--r-- 1 root hadoop 110990 2026-09-02 08:36 /churn_proj/customers/204
088a4-105a-4420-8931-ad2ec14b170e
-rw-r--r-- 1 root hadoop 110088 2026-09-02 08:36 /churn_proj/customers/444
74738-d385-4864-99c3-897cc3170cac
-rw-r--r-- 1 root hadoop 111055 2026-09-02 08:36 /churn_proj/customers/5f1
8d9b6-2400-4af3-bb7a-ad1b123ca063
-rw-r--r-- 1 root hadoop 111021 2026-09-02 08:36 /churn_proj/customers/677
d9eb9-0e52-4bff-80f2-a2eff4a48a9b
-rw-r--r-- 1 root hadoop 111134 2026-09-02 08:36 /churn_proj/customers/713
8bda3-4a4c-4cda-8b8d-5c0dd7198730
-rw-r--r-- 1 root hadoop 111162 2026-09-02 08:36 /churn_proj/customers/a60
430cc-3c5a-4e0a-9d9d-a8e8e3293b52
-rw-r--r-- 1 root hadoop 110982 2026-09-02 08:36 /churn_proj/customers/e37
12207-de5a-425c-a4c7-9abef87fbbcd
-rw-r--r-- 1 root hadoop 111121 2026-09-02 08:36 /churn_proj/customers/eab
dc02e-de77-4fb3-906b-14aa7d9ff568
-rw-r--r-- 1 root hadoop 111131 2026-09-02 08:36 /churn_proj/customers/f1b
9c7c4-812d-4b47-9e4d-8aefe6ac6093


```

```

[student@localhost ~]$ hdfs dfs -cat /churn_proj/customers/0f3eee39-977c-421b-9c
f9-c85479f78b8d | head -n 5
row_id,customer_id,age,gender,geography,is active,tenure_months,date opened,cred
it score,balance,num_products,has_cr_card,estimated_salary,is churned,source_fil
e,load_batch ts
2001,15613656,58.0,Male,France,1,12,2025-09-01,842,63492.94,1,1,83172.19,0,custo
mers.csv,2026-09-02 02:28:03.0
2002,15734311,27.0,Female,France,1,36,2023-09-02,661,0.0,2,1,76889.79,0,customer
s.csv,2026-09-02 02:28:03.0
2003,15657214,74.0,Male,France,1,24,2024-09-01,601,0.0,2,0,51554.58,0,customer
s.csv,2026-09-02 02:28:03.0
2004,15799350,41.0,Male,france,1,0,2026-09-01,632,106134.46,1,0,105570.39,0,cust
omers.csv,2026-09-02 02:28:03.0

```

- **Ingestion Status:** 11 non-empty CSV files (~111 KB each) were successfully generated by NiFi and stored in /churn_proj/customers.
- **File Permissions:** All dataset files are owned by root:hadoop with standard -rw-r--r-- read/write access rights.
- **Validation Command:** Executed `hdfs dfs -cat /churn_proj/customers/<file_id> | head -n 5` to inspect schema integrity and record formatting.
- **Header Structure:** Confirmed the presence of comma-delimited columns (e.g., customer_id, first_name, last_name, email, gender, churn_status).
- **Verification Result:** The output confirms valid CSV formatting, correct column headers, and successful incremental data landing in HDFS.



2. Other Files ingestion

A lightweight ingestion pipeline used to fetch raw file sources from the local filesystem and store them directly into HDFS.

Processed Datasets

- **CSV Datasets:** customer_support_tickets.csv, offers.csv
- **JSON Dataset:** usage.json

Component Summary

- **GetFile:** Polls a specified local directory to ingest raw CSV and JSON files into NiFi FlowFiles upon detection.
- **PutHDFS:** Writes the ingested files directly into their designated target directories on the HDFS cluster.

GetFile Processor

Configure Processor

Stopped

SETTINGS SCHEDULING PROPERTIES COMMENTS

Required field

Property	Value
Input Directory	/media/sf_Project
File Filter	*usage*.json\$
Path Filter	No value set
Batch Size	10
Keep Source File	true
Recurse Subdirectories	false
Polling Interval	0 sec
Ignore Hidden Files	true
Minimum File Age	0 sec
Maximum File Age	No value set
Minimum File Size	0 B
Maximum File Size	No value set


CANCEL APPLY

- **Input Directory (/media/sf_Project):** Specifies the shared local directory where the raw source files are picked up.
- **File Filter:** Uses regular expressions to target and ingest specific files individually:
 - ^usage\.json\$ (for the JSON file)
 - ^customer_support_tickets\.csv\$ (for the tickets CSV file)
 - ^offers\.csv\$ (for the offers CSV file)
- **Keep Source File (true):** Retains the original raw files in the source directory after ingestion instead of deleting.
- **Batch Size (10):** Limits processing to a maximum of 10 files per execution cycle.

```
[student@localhost ~]$ hdfs dfs -ls /churn_proj/usage
Found 1 items
-rw-r--r--  1 root hadoop   18363496 2026-09-02 09:48 /churn_proj/usage/usage.json


[student@localhost ~]$ hdfs dfs -ls /churn_proj/customer_support_tickets
Found 1 items
-rw-r--r--  1 root hadoop    508525 2026-09-02 09:26 /churn_proj/customer_support_tickets/customer_support_tickets.csv

[student@localhost ~]$ hdfs dfs -ls /churn_proj/offers
Found 1 items
-rw-r--r--  1 root hadoop    358194 2026-09-02 09:39 /churn_proj/offers/offers.csv
```



2. PySpark Data Processing Pipeline (Medallion Architecture)

The PySpark ETL pipeline executes on a distributed cluster to process raw data landed in HDFS through Bronze, Silver, and Gold transformation stages.



2.1 Bronze Layer (Raw Ingestion)

The Bronze layer contains the raw, uncleaned data as received from the source files. At this stage, the data is loaded into PySpark DataFrames without applying any cleaning or transformation operations.

For each of the four source tables, the schema is printed to display the available features and their original data types. The number of rows is also calculated and reported to provide an initial overview of the raw datasets.

No data cleaning, type conversion, or feature modification is performed in the Bronze layer. The purpose of this layer is to preserve and inspect the data in its original form before applying transformations in the subsequent layers.

3 BRONZE LAYER — Raw Ingestion from HDFS

Straight reads of what NiFi landed in `/churn_proj/...`. No cleaning here — this is the raw, as-received data, kept exactly as ingested.

Bronze — Customers

```
In [3]: df_customers = (
        spark.read
        .option("header", True)
        .csv(f"{BASE}/customers")
        )

print("Rows:", df_customers.count())
df_customers.printSchema()

Rows: 10150
root
 |-- row_id: string (nullable = true)
 |-- customer_id: string (nullable = true)
 |-- age: string (nullable = true)
 |-- gender: string (nullable = true)
 |-- geography: string (nullable = true)
 |-- is_active: string (nullable = true)
 |-- tenure_months: string (nullable = true)
 |-- date_opened: string (nullable = true)
 |-- credit_score: string (nullable = true)
 |-- balance: string (nullable = true)
 |-- num_products: string (nullable = true)
 |-- has_cr_card: string (nullable = true)
 |-- estimated_salary: string (nullable = true)
 |-- is_churned: string (nullable = true)
 |-- source_file: string (nullable = true)
 |-- load_batch_ts: string (nullable = true)
```

2.2 Silver Layer (Cleaning and Transformation)

The Silver layer is the **cleaning and conformance layer** between raw Bronze data and the downstream warehouse/Gold layer.

The processing implemented in the notebook performs:

1. Data-type standardization
2. Whitespace trimming
3. Sentinel-value normalization
4. Categorical-value standardization
5. Domain validation
6. Duplicate detection and removal
7. Referential-integrity checks against `customer_id`
8. Missing-value analysis
9. Missing-value imputation
10. Final schema and NULL validation

2. Silver — Customer Features

2.1 Source

The Bronze customer source initially contained **10,150 rows**, with all business columns represented as strings.

Bronze — Customers

```
df_customers = (
    spark.read
    .option("header", True)
    .csv(f"{BASE}/customers")
)

print("Rows:", df_customers.count())
df_customers.printSchema()

Rows: 10150
root
 |-- row_id: string (nullable = true)
 |-- customer_id: string (nullable = true)
 |-- age: string (nullable = true)
 |-- gender: string (nullable = true)
 |-- geography: string (nullable = true)
 |-- is_active: string (nullable = true)
 |-- tenure_months: string (nullable = true)
 |-- date_opened: string (nullable = true)
 |-- credit_score: string (nullable = true)
 |-- balance: string (nullable = true)
 |-- num_products: string (nullable = true)
 |-- has_cr_card: string (nullable = true)
 |-- estimated_salary: string (nullable = true)
 |-- is_churned: string (nullable = true)
 |-- source_file: string (nullable = true)
 |-- load_batch_ts: string (nullable = true)
```

Initial problem

The notebook explicitly notes that the customer features initially had **string data types**, even for numeric, Boolean, date, and timestamp attributes.

2.2 General preprocessing

`row_id`  **Removed**

`row_id` was treated as a technical identifier rather than a business attribute and was removed from the customer dataset.

Reason: It does not represent a meaningful customer business attribute and is unnecessary in the cleaned customer dataset.

row_id — Removed

row_id was only used as a technical identifier and is not needed for the final customer dataset, so it was removed.

```
# Trim whitespace from all columns (Now All string DataType)

cols = df_customers.columns

for c in cols:
    df_customers = df_customers.withColumn(
        c,
        F.trim(F.col(c))
    )
```

null , N/A , unknown Values require to be unified

```
# Convert sentinel values to NULL
for c in cols:
    df_customers = df_customers.withColumn(
        c,
        F.when(
            F.upper(F.trim(F.col(c))).isin(
                "N/A",
                "UNKNOWN",
                "NULL"
            ),
            None
        )
        .otherwise(F.col(c))
    )
```

White space

White space was trimmed from **all customer columns** before further processing.

Transformation: TRIM(column)

Purpose: Prevent values such as " France " or " Male " from being treated as different values from "France" or "Male".

Sentinel values

The notebook identifies N/A, UNKNOWN, and NULL string representations and converts them into actual database NULL values.

This provides a consistent representation of missing/unknown data.

3. Customer Feature-by-Feature preprocessing

customer_id

Role: Customer business identifier.

Silver transformation: No value transformation is documented for customer_id.

Validation: It is retained as a string and subsequently used as the reference key for validating Offers and Tickets.

age

Problems found

The source contained:

- normal numeric values such as 25, 37, 58
- values represented as 37.0
- textual values such as "years 42"
- invalid values such as -1, -5
- extreme values such as 999, 9999, 121, 150
- NULLs

The EDA identified **82 NULLs initially**.

Transformation

The notebook extracts the numeric portion:

"years 42" → 42

"25" → 25

and converts the result to INTEGER.

Then domain validation is applied: **age < 15 OR age > 100 → NULL**

```
df_customers = df_customers.withColumn(
    "age",
    F.when(
        (F.col("age") < 15) | (F.col("age") > 100),
        None
    ).otherwise(F.col("age"))
)
```

This converted clearly invalid ages into NULL. After this validation, the NULL count became **128**.

```
: df_customers.filter(
    F.col("age").isNull()
).count()
```

: 128

Missing-value treatment

An IQR analysis produced:

- Q1 = 32
- Q3 = 44
- Median = 37
- Lower bound = 14
- Upper bound = 62

The NULL ages were subsequently filled using the median **37**.

Final

age → INTEGER

invalid/missing age → median 37



IQR for Detecting Outliers

```
def IQR(df, column):
    q1, q3 = df.approxQuantile(column, [0.25, 0.75], 0.01)
    median = df.approxQuantile(column, [0.5], 0.01)[0]

    iqr = q3 - q1
    lower = q1 - 1.5 * iqr
    upper = q3 + 1.5 * iqr

    print(f"{column}")
    print(f"Q1 : {q1}")
    print(f"Q3 : {q3}")
    print(f"Median : {median}")
    print(f"IQR : {iqr}")
    print(f"Lower : {lower}")
    print(f"Upper : {upper}")

    return median
```

```
age_median = IQR(df_customers, "age")
```

```
age
Q1   : 32.0
Q3   : 44.0
Median : 37.0
IQR   : 12.0
Lower : 14.0
Upper : 62.0
```

```
age NULL ==>> Median
```

gender

Problems found

The source contained:

- Male
- Female
- male
- FEMALE
- NULL

The EDA showed:

Value	Count
Male	5,393
Female	4,488
male	45
NULL	37
FEMALE	37

```
df_customers.groupBy("gender") \
  .count() \
  .orderBy(F.desc("count")) \
  .show(100, truncate=False)
```

[Stage 61:=====> (169 + 7) / 200]

gender	count
Male	5393
Female	4488
male	45
FEMALE	37
null	37

Transformation

Values are standardized:

MALE / M → Male
FEMALE / F → Female

The implemented transformation normalizes the values using uppercase comparisons.

Missing-value treatment

The mode was calculated as: Male

NULL gender values were filled with the mode.

Final

gender → Male / Female
NULL → Male

The final Silver schema marks gender as non-nullable.

Geography

Problems found

The source contained:

- France
- Germany
- Spain
- france

- GERMANY
- ES
- NULL

EDA showed the inconsistent representations.

Transformation

The notebook maps:

FRANCE / FR → France
 GERMANY / DE → Germany
 SPAIN / ES → Spain

Missing-value treatment

The mode was: France

NULL geography values were replaced with France.

Final

geography → France / Germany / Spain
 NULL → France

```
df_customers.groupBy("geography") \
  .count() \
  .orderBy(F.desc("count")) \
  .show(100, truncate=False)
```

[Stage 58:=====]

geography	count
France	4940
Germany	2475
Spain	2437
null	80
france	31
ES	22
GERMANY	15

```
# Normalize geography
df_customers = df_customers.withColumn(
  "geography",
  F.when(F.upper(F.col("geography")).isin("FRANCE", "FR"), "France")
  .when(F.upper(F.col("geography")).isin("GERMANY", "DE"), "Germany")
  .when(F.upper(F.col("geography")).isin("SPAIN", "ES"), "Spain")
  .otherwise(F.col("geography"))
)
```

is_active and is_churned

Problems found

The source values were: 0 & 1

with:

- 5,151 active
- 4,849 inactive
- 0 NULLs

Transformation

Converted from string/integer representation into Boolean:

1 → True && 0 → False

Final : is_active & is_churned → BOOLEAN

is_active and is_churned

```
df_customers.groupBy("is_active") \
  .count() \
  .orderBy(F.desc("count")) \
  .show()
df_customers.groupBy("is_churned") \
  .count() \
  .orderBy(F.desc("count")) \
  .show()
```

is_active	count
1	5151
0	4849

[Stage 73:=====] (171 + 6) / 200

is_churned	count
0	7963
1	2037

```
print(df_customers.filter(F.col("is_active").isNull()).count())
print(df_customers.filter(F.col("is_churned").isNull()).count())
```

0
0

date_opened

Problems found

The column contained many different date representations, including:

2026-06-13
 04/03/2026
 07-03-2026

May 31, 2026
2026/06/13

There were initially **0 NULL dates**.

Transformation

Multiple formats are attempted using coalesce:

yyyy-MM-dd
MMM dd, yyyy
dd/MM/yyyy
dd-MM-yyyy
yyyy/MM/dd
MM-dd-yyyy

Final

date_opened → DATE

This standardizes all valid representations to a single Spark DATE type.

```
In [16]: # normalize the dates
df_customers = df_customers.withColumn(
    "date_opened",
    F.coalesce(
        F.to_date("date_opened", "yyyy-MM-dd"),
        F.to_date("date_opened", "MMM dd, yyyy"),
        F.to_date("date_opened", "dd/MM/yyyy"),
        F.to_date("date_opened", "dd-MM-yyyy"),
        F.to_date("date_opened", "yyyy/MM/dd"),
        F.to_date("date_opened", "MM-dd-yyyy")
    )
)
```

```
In [17]: df_customers.select("date_opened") \
    .show(100, truncate=False)

df_customers.printSchema()

+-----+
|2019-09-03|
|2021-09-02|
|2020-09-02|
|2017-09-03|
|2018-09-03|
|2019-09-03|
|2020-09-02|
|2021-09-02|
|2017-09-03|
|2021-09-02|
|2022-09-02|
|2024-09-01|
|2019-09-03|
|2017-09-03|
|2017-09-03|
|2017-09-03|
|2018-09-03|
|2025-09-01|
```

tenure_months

Problems found

Valid values include:

0, 12, 24, 36, 48, 60,
72, 84, 96, 108, 120

```
|108|980|
|72|962|
|120|488|
|0|412|
|999|13|
|-1|11|
|-10|10|
+-----+-----+
```

Invalid values include:

999
-1
-10

```
df_customers = df_customers.withColumn(
    "tenure_months",
    F.when(
        (F.col("tenure_months") < 0) |
        (F.col("tenure_months") == 999),
        None
    ).otherwise(F.col("tenure_months").cast("int"))
)
```

Transformation

negative values → NULL
999 → NULL
otherwise → INTEGER

Final

tenure_months → INTEGER

Valid tenure values are preserved.

tenure_months — EDA Findings

The `tenure_months` column was stored as a `STRING`.

The column contains valid tenure values such as 0, 12, 24, 36, 48, 60, 72, 84, 96, 108, 120 months.

Some invalid values were also found:

- 999 → invalid/corrupted value
- -1 → invalid negative value
- -10 → invalid negative value

Cleaning applied:

- Converted `tenure_months` from `STRING` to `INTEGER`.
- Replaced 999 and negative values with `NULL`.
- Preserved all valid tenure values.

credit_score

Problems found

Most values were numeric, but corrupted representations were found:

"six hundred"
"score_700"
"unknown"

Transformation

"six hundred" → 600
"score_700" → 700

and the remaining values are cast to integer.

Missing-value treatment

IQR analysis:

- Q1 = 582
- Q3 = 715
- Median = **650**
- Lower bound = 382.5
- Upper bound = 914.5

Missing credit scores were filled with **650**.

Final

credit_score → INTEGER
NULL → 650

```
df_customers.filter(
  F.lower(F.col("credit_score")).isin(
    "six hundred",
    "score_700",
    "unknown"
  )
).groupBy("credit_score").count().show()

+-----+
|credit_score|count|
+-----+
|   score_700|   23|
|  six hundred|   18|
+-----+
```

credit_score — EDA Findings

The column contains corrupted categorical representations:

- "six hundred" → should be mapped to 600
- "score_700" → should be mapped to 700
- "unknown" was mapped to NULL at first

```
df_customers = df_customers.withColumn(
  "credit_score",
  F.when(F.col("credit_score") == "six hundred", 600)
  .when(F.col("credit_score") == "score_700", 700)
  .otherwise(F.col("credit_score").cast("int"))
)
df_customers.groupBy("credit_score").count().show(20)

[Stage 109:=====> (188 + 6) / 200]

+-----+
|credit_score|count|
+-----+
|          471|    7|
+-----+
```

balance

Problems found

The EDA identified:

- 0.0 → valid
- NULL → missing
- -1, -1.0 → invalid sentinel
- -5000.0 → invalid/suspicious
- 9999 → suspicious
- 9999999.0 → extreme suspicious value

Transformation

The following values were converted to NULL:

-1
-1.0
-5000.0
9999
9999999.0

```
df_customers = df_customers.withColumn(
  "balance",
  F.when(
    F.col("balance").isin("-1", "-1.0", "-5000.0", "9999", "9999999.0"),
    None
  )
  .otherwise(F.col("balance"))
)
```

```
df_customers.select(
  F.min("balance").alias("min"),
  F.max("balance").alias("max"),
  F.avg("balance").alias("avg")
).show()

[Stage 136:=====> (148 + 6) / 200]

+-----+
|min|   max|   avg|
+-----+
|0.0|99986.98|76472.81348644278|
+-----+
```

```
df_customers = df_customers.withColumn(
  "balance",
  F.col("balance").cast("double")
)
```

```
df_customers.groupBy("balance").count().orderBy(F.desc("count")).show(20)

[Stage 139:=====> (183 + 6) / 200]

+-----+
| balance|count|
+-----+
|      0.0|  3575|
|      null|   116|
+-----+
```

Then the column was cast to: DOUBLE

After cleaning the range was:

min = 0
max = 99,986.98
avg $\approx$ 76,472.81

Missing-value treatment

Median was calculated as approximately **96,402.96**, and NULL balances were replaced by the median.

Final

balance $\rightarrow$ DOUBLE
NULL/invalid $\rightarrow$ median

num_products

Problems found

Valid values: 1, 2, 3, 4

Corrupted/suspicious values:

"three" "two" "many" -1 0 7 99

Transformation

"three" $\rightarrow$ 3
"two" $\rightarrow$ 2
"many" $\rightarrow$ NULL

Then values are converted to integers.

Domain validation then applies:

value < 1 OR value > 4 $\rightarrow$ NULL

After validation:

Value Count

1	5,049
2	4,567
3	283
4	60
NULL	41

Missing-value treatment

The mode was **1**, and missing values were filled with 1

Final

num_products $\rightarrow$ INTEGER
invalid $\rightarrow$ NULL $\rightarrow$ mode 1

num_products — EDA Findings

The column contains both valid numeric values and corrupted representations:

- 1, 2, 3, 4 $\rightarrow$ valid values From data description
- "three" $\rightarrow$ should be mapped to 3
- "two" $\rightarrow$ should be mapped to 2
- "many" $\rightarrow$ invalid/unknown $\rightarrow$ NULL
- 7, 99 $\rightarrow$ suspicious values that need validation
- 0, -1 $\rightarrow$ invalid values

Cleaning required:

1. Convert "three" $\rightarrow$ 3
2. Convert "two" $\rightarrow$ 2
3. Convert "many" $\rightarrow$ NULL
4. Handle invalid/suspicious numeric values (-1, 0, 7, 99)
5. Convert the column to integer

```
# mapping values
df_customers = df_customers.withColumn(
    "num_products",
    F.when(F.col("num_products") == "three", 3)
    .when(F.col("num_products") == "two", 2)
    .when(F.col("num_products") == "many", None)
    .otherwise(F.col("num_products").cast("int"))
)
df_customers.groupBy("num_products").count().orderBy("num_products").show()

[Stage 85:=====> (157 + 6) / 200]

+-----+
|num_products|count|
+-----+
|          null|   13|
|          -1|    9|
|           0|    6|
+-----+
```

```
num_products_mode = (
    df_customers
    .filter(F.col("num_products").isNotNull())
    .groupBy("num_products")
    .count()
    .orderBy(F.desc("count"))
    .first()["num_products"]
)

print("num_products mode:", num_products_mode)

df_customers = df_customers.fillna({
    "num_products": num_products_mode
})

[Stage 214:=====> (155 + 6) / 200]

num_products mode: 1
```

has_cr_card

Source

Values were:

0 → 2,945

1 → 7,055

with no NULL values shown.

Transformation

Converted to Boolean.

0 → False

1 → True

Final

has_cr_card → BOOLEAN

has_cr_card

```
df_customers.groupBy("has_cr_card") \
    .count() \
    .orderBy("has_cr_card") \
    .show()
```

[Stage 142:=====> (163 + 6) / 200]

```
+-----+-----+
|has_cr_card|count|
+-----+-----+
|          0| 2945|
|          1| 7055|
+-----+-----+
```

```
df_customers = df_customers.withColumn("has_cr_card",F.col("has_cr_card").cast("boolean"))
```

estimated_salary

EDA

The notebook checks the distribution and confirms **0 NULL values**.

Transformation

Converted from string to => DOUBLE

Final

estimated_salary → DOUBLE

No imputation was required.

```
: df_customers.filter(
:     F.col("estimated_salary").isNull()
: ).count()
```

: 0

```
: df_customers = df_customers.withColumn(
:     "estimated_salary",
:     F.col("estimated_salary").cast("double")
: )
```

```
: df_customers.printSchema()
```

```
root
|-- customer_id: string (nullable = true)
|-- age: integer (nullable = true)
|-- gender: string (nullable = true)
|-- geography: string (nullable = true)
|-- is_active: boolean (nullable = true)
|-- tenure_months: string (nullable = true)
|-- date_opened: date (nullable = true)
|-- credit_score: integer (nullable = true)
|-- balance: double (nullable = true)
|-- num_products: integer (nullable = true)
|-- has_cr_card: boolean (nullable = true)
|-- estimated_salary: double (nullable = true)
|-- is_churned: boolean (nullable = true)
|-- source_file: string (nullable = true)
|-- load_batch_ts: timestamp (nullable = true)
```

4. Customer Duplicate Handling

We found **150 duplicate groups** based on the business/customer attributes.

The dataset was then deduplicated using: dropDuplicates()

Result : Duplicate customer records were removed before the downstream Silver output.

```
: check_cols = [
:     c for c in df_customers.columns
:     if c not in ["row_id", "source_file", "load_batch_ts"]
: ]
```

```
df_customers.groupBy(check_cols) \
    .count() \
    .filter(F.col("count") > 1) \
    .count()
```

: 150

```
: df_customers = df_customers.dropDuplicates()
```

5. Customer Missing-Value Strategy

The final customer imputation strategy is:

Feature	Missing-value treatment
gender	Mode = Male
geography	Mode = France
age	Median = 37
credit_score	Median = 650
balance	Median $\approx$ 96,402.96
num_products	Mode = 1

Silver — Offers

The Offers source contains the following columns:

- offer_id
- customer_id
- offer_type
- accepted
- date_offered

The Bronze dataset contained **8,127 records**.

Bronze — Offers

```
df_offers = spark.read.option("header", True).csv(f"{BASE}/offers")
print("Raw rows:", df_offers.count())
df_offers.printSchema()
```

```
Raw rows: 8127
root
|-- offer_id: string (nullable = true)
|-- customer_id: string (nullable = true)
|-- offer_type: string (nullable = true)
|-- accepted: string (nullable = true)
|-- date_offered: string (nullable = true)
```

offer_id

Role: Offer identifier.

No value transformation is applied to the offer_id itself.

Duplicate Handling

The notebook identified **107 duplicate offer_id groups**.

Duplicate offers were removed using the following operation:

dropDuplicates(["offer_id"])

NULL offer_id values were not removed by this operation and were explicitly preserved during the validation process.

customer_id

The Offer customer_id was validated against the cleaned Customer dataset.

The notebook performed a left-anti join to identify customer IDs that do not exist in the Customer dataset.

The validation identified **144 orphan customer_id rows**.

These orphan records were then removed using a left_semi join against the valid customer IDs.

Result

Only Offer records whose customer_id exists in the cleaned Customer dataset remain in the Silver layer.

offer_type

Problems Identified

The source contained the following expected offer types:

- Fee Waiver
- Cashback
- Loyalty Bonus
- Product Upgrade
- Rate Discount

Several casing inconsistencies were also found, including:

- fee waiver
- cashback
- product upgrade
- RATE DISCOUNT

The dataset also contained missing or sentinel representations such as:

- Unknown
- N/A
- NULL
- actual NULL

Transformation

The casing of the offer types was standardized:

- fee waiver → Fee Waiver
- cashback → Cashback
- product upgrade → Product Upgrade
- rate discount → Rate Discount

The values Unknown, N/A, and string representations of NULL were converted into actual NULL values.

Imputation

After standardization, there were **78 NULL values** in offer_type.

These missing values were replaced with:

Unknown

Therefore, the final Silver dataset does not contain missing offer_type values.

accepted

Problems Identified

The source contained several different representations of acceptance status:

- True
- False
- TRUE
- TRUEEE
- 1
- 0
- yes
- no
- MAYBE
- 2
- NULL

Transformation

Values representing an accepted offer were standardized as follows:

- TRUE → True
- TRUEEE → True
- 1 → True
- YES → True

Values representing a rejected offer were standardized as follows:

- FALSE → False
- 0 → False
- NO → False

Any other value that could not be reliably interpreted as accepted or rejected was converted to NULL.

Important Design Decision

The notebook does **not** convert unknown acceptance values into False.

Instead, the original accepted column remains a nullable Boolean field:

accepted → BOOLEAN (nullable)

A separate accepted_flag field is also created to provide a user-friendly categorical representation:

- True → Yes
- False → No
- NULL → Unknown

This approach prevents unknown acceptance statuses from being incorrectly interpreted as rejected offers.

date_offered

Problems Identified

Multiple date formats were found in the source data, including:

- yyyy-MM-dd
- MM-dd-yyyy
- yyyy/MM/dd
- dd/MM/yyyy
- dd-MM-yyyy
- MMM dd, yyyy

Transformation

The different supported date formats were parsed and standardized into the Spark DATE data type.

Therefore:

date_offered → DATE

Missing Values

There were **64 NULL values** in date_offered.

The original date_offered column remains nullable because a missing business date should not be replaced with a fabricated real-world date.

A separate field called:

date_offered_key

is created for downstream date-dimension processing.

For missing dates:

NULL → 1900-01-01

The value 1900-01-01 is used as a technical/default date key so that records with missing offer dates can still be joined cleanly to the Date dimension without claiming that the offer actually occurred on that date.

7. Silver — Tickets

The Ticket source contains the following columns:

- ticket_id
- customer_id
- issue_type
- severity
- resolution_time_hrs

The source contains **14,182 raw records**.

ticket_id

Role: Ticket identifier.

Bronze — Tickets

```
df_tickets = spark.read.option("header", True).csv(f"{BASE}/customer_support_tickets")
print("Raw rows:", df_tickets.count())
df_tickets.printSchema()

Raw rows: 14182
root
 |-- ticket_id: string (nullable = true)
 |-- customer_id: string (nullable = true)
 |-- issue_type: string (nullable = true)
 |-- severity: string (nullable = true)
 |-- resolution_time_hrs: string (nullable = true)
```

No transformation is applied to the identifier itself.

Duplicate records were checked across the complete record. The notebook found **210 duplicate groups**, and exact duplicate records were removed.

customer_id

Ticket customer IDs were checked against the valid customers in the cleaned Customer dataset.

The notebook found **209 non-null orphan customer IDs**.

These orphan records were removed using a left_semi join against the valid customer IDs.

Final Validation

After the cleaning process: **customer_id NULLs = 0**

This ensures that every remaining ticket is associated with a valid customer.

issue_type

Problems Identified

The expected issue types are:

- App Bug
- Loan Query
- Billing
- Fraud Alert
- Card Problem
- Login Issue

Several inconsistent representations were found, including:

- login issue
- app bug
- loan query
- billing
- FRAUD ALERT
- CARD PROBLEM

The dataset also contained missing-value representations such as:

- Unknown
- N/A
- NULL
- actual NULL

Transformation

The different casing representations were standardized to the expected values:

- app bug → App Bug
- loan query → Loan Query
- billing → Billing
- fraud alert → Fraud Alert

- card problem → Card Problem
- login issue → Login Issue

The values Unknown, N/A, and string representations of NULL were converted to actual NULL values.

Imputation

After standardization, the NULL values were handled using the most frequent value (mode).

The calculated mode was:

Loan Query

Therefore, missing issue_type values were replaced with **Loan Query**.

severity

Problems Identified

The expected severity values are:

- Low
- Medium
- High
- Critical

Inconsistent representations were also found, including:

- low
- MEDIUM
- high
- CRITICAL

The dataset also contained missing-value representations such as:

- Unknown
- N/A
- NULL

Transformation

All valid severity values were standardized to the following format:

- Low
- Medium
- High
- Critical

The values Unknown, N/A, and string representations of NULL were converted to actual NULL values.

Imputation

The most frequent severity value (mode) was: **Low**

Therefore, missing severity values were replaced with **Low**.

resolution_time_hrs

Problems Identified

The source contained:

- Valid decimal resolution times
- NULL values
- Negative values such as -1 and -5
- Sentinel values such as 999, 9999, and 2500

Transformation

The original field was stored as a string and was converted to a numeric data type.

STRING → DOUBLE

Invalid and sentinel values were converted to NULL:

- Negative values → NULL
- 999 → NULL
- 9999 → NULL
- 2500 → NULL

After the cleaning process, the observed range was:

- Minimum = **1.0**
- Maximum = **119.9**
- NULL values = **449**

Imputation

An IQR-based analysis was performed to determine the appropriate value for missing resolution times.

The calculated statistics were:

- Q1 = **2.4**
- Q3 = **15.3**
- Median = **6.5**
- IQR = **12.9**

The **449 missing values** were replaced with the median value:

6.5 hours

Final State

- resolution_time_hrs → **DOUBLE**
 - NULL values → **6.5**
-

8. Silver — Usage

The Usage dataset contains **120,000 records**.

The dataset contains the following columns:

- customer_id
- log_date
- monthly_balance
- num_products
- product_type
- usage_log_id

Missing Values

A complete NULL audit was performed across all Usage columns.

The result was:

0 NULL values

Therefore, no missing-value imputation or record removal was required for the Usage dataset.

This is different from the Customers, Offers, and Tickets datasets, where missing and invalid values required additional treatment.

Silver Treatment

Based on the Silver-layer implementation in the notebook, the Usage dataset was validated for completeness, but no additional detailed cleaning or type-conformance transformations were implemented.

The existing data types remain:

Column	Data Type
customer_id	LONG
log_date	STRING
monthly_balance	DOUBLE
num_products	LONG
product_type	STRING
usage_log_id	LONG

Therefore, the Usage Silver-layer treatment can be summarized as:

Validated for completeness, with no NULL remediation required. Existing data types were retained in the current notebook implementation.

3. Gold Layer - Data Warehouse Architecture (Apache Hive)

This section describes the technical implementation of the **Gold Layer** within the Medallion Data Architecture. The primary objective is to transform cleansed, columnar data from the **Silver Layer** into a business-ready **Star Schema** optimized for Business Intelligence (BI) and analytical workloads. The pipeline was implemented using **Apache Hive** within the Hadoop ecosystem.

2. Environment & Session Configurations

To ensure stable execution and allow for advanced table partitioning within Hive, specific session parameters and execution engines were configured before running the ETL jobs.


SQL

```
-- Enable dynamic partitioning for the Fact table
```

```
1 SET hive.exec.dynamic.partition = true;
2 SET hive.exec.dynamic.partition.mode = nonstrict;
3
4
```

3. Data Definition Language (DDL) - Star Schema Design

The data model was structured into one partitioned Fact table and multiple Dimension tables. All tables were created as EXTERNAL tables stored in PARQUET format to decouple compute from storage and optimize read performance.



SQL

```
4
5 CREATE DATABASE IF NOT EXISTS project1_dw;
6 USE project1_dw;
7
```

```
-- =====
```

```
-- Dimension Tables
```

```
-- =====
```

```
-- 1. Dim_Customer
```

```

73
74 -- Dim_Customer
75 CREATE EXTERNAL TABLE IF NOT EXISTS Dim_Customer (
76     cust_key INT,
77     customer_id BIGINT,
78     age INT,
79     gender VARCHAR(20),
80     geography VARCHAR(50),
81     credit_score INT,
82     balance DECIMAL(15,2),
83     tenure_months INT,
84     num_products INT,
85     has_cr_card BOOLEAN,
86     is_active BOOLEAN,
87     estimated_salary DECIMAL(15,2),
88     is_churned BOOLEAN,
89     CLV_LTV DECIMAL(15,2),
90     avg_ticket_res_time_hrs DECIMAL(10,2),
91     dw_start_date DATE,
92     dw_end_date DATE
93 ) STORED AS PARQUET
94 LOCATION '/user/student/project1/gold/parquet/Dim_Customer';
95
-- 1. Dim_Product
114
115 CREATE EXTERNAL TABLE IF NOT EXISTS Dim_Product (
116     product_key INT,
117     product_type VARCHAR(100)
118 ) STORED AS PARQUET
119 LOCATION '/user/student/project1/gold/parquet/Dim_Product';
120
-- 2. Dim_Offer
97
98 CREATE EXTERNAL TABLE IF NOT EXISTS Dim_Offer (
99     offer_key INT,
100     offer_type VARCHAR(100)
101 ) STORED AS PARQUET
102 LOCATION '/user/student/project1/gold/parquet/Dim_Offer';
103
-- 3. Dim_Ticket
103
104 -- Dim_Ticket (Category Based)
105
106 CREATE EXTERNAL TABLE IF NOT EXISTS Dim_Ticket (
107     ticket_key INT,
108     issue_type VARCHAR(100),
109     severity VARCHAR(50)
110 ) STORED AS PARQUET
111 LOCATION '/user/student/project1/gold/parquet/Dim_Ticket';
112
-- 4. Dim_Date
102
103 -- Dim_Date
104 CREATE EXTERNAL TABLE IF NOT EXISTS Dim_Date (
105     date_key INT,
106     full_date DATE,
107     day INT,
108     month INT,
109     month_name VARCHAR(20),
110     quarter INT,
111     year INT
112 ) STORED AS PARQUET
113 LOCATION '/user/student/project1/gold/parquet/Dim_Date';
114
-- 5. Fact_Customer_Activity (Dynamic Partitioning)
136
137 CREATE EXTERNAL TABLE IF NOT EXISTS Fact_Customer_Activity (
138     activity_key BIGINT,
139     cust_key INT,
140     date_key INT,
141     ticket_key INT,
142     offer_key INT,
143     product_key INT,
144     usage_balance DECIMAL(15,2),
145     resolution_time_hrs DECIMAL(10,2),
146     is_accepted BOOLEAN
147 )
148 PARTITIONED BY (event_type VARCHAR(20), event_year INT)
149 STORED AS PARQUET
150 LOCATION '/user/student/project1/gold/parquet/Fact_Customer_Activity';
151

```

4. Data Manipulation Language (DML) - ETL Processing

The ETL process involved complex transformations, including the programmatic calculation of Customer Lifetime Value (CLV/LTV), hashing for surrogate key generation (ABS(HASH())), and handling orphan records using sentinel values (e.g., -1).

SQL

```
-- =====
-- Loading Dimensions
-- =====

-- Load Dim_Customer (with programmatic CLV_LTV calculation)
WITH Customer_Ticket_Avg AS (
  SELECT customer_id, AVG(resolution_time_hrs) AS avg_res
  FROM tickets
  GROUP BY customer_id
)
INSERT OVERWRITE TABLE Dim_Customer
SELECT
  ABS(HASH(c.customer_id)) AS cust_key,
  CAST(c.customer_id AS BIGINT),
  c.age, c.gender, c.geography, c.credit_score,
  CAST(c.balance AS DECIMAL(15,2)), c.tenure_months, c.num_products,
  c.has_cr_card, c.is_active, CAST(c.estimated_salary AS DECIMAL(15,2)), c.is_churned,

  -- CLV / LTV Logic
  CAST(
    CASE
      WHEN c.is_churned = TRUE THEN 0
      ELSE c.balance * c.num_products * (
        CASE
          WHEN c.tenure_months < 12 THEN 5
          WHEN c.tenure_months < 36 THEN 4
          WHEN c.tenure_months < 60 THEN 3
          ELSE 2
        )
      END
    ) AS DECIMAL(15,2)
  ) AS CLV_LTV,

  CAST(t.avg_res AS DECIMAL(10,2)) AS avg_ticket_res_time_hrs,
  c.date_opened AS dw_start_date,
  NULL AS dw_end_date
FROM customers c
LEFT JOIN Customer_Ticket_Avg t ON c.customer_id = t.customer_id;

-- Load Dim_Offer
212
213 INSERT OVERWRITE TABLE Dim_Offer
214 SELECT DISTINCT ABS(HASH(offer_type)) AS offer_key, offer_type
215 FROM offers WHERE offer_type IS NOT NULL;
216
217 INSERT INTO TABLE Dim_Offer VALUES (-1, 'No Offer');
218

-- Load Dim_Ticket
220
221 INSERT OVERWRITE TABLE Dim_Ticket
222 SELECT DISTINCT ABS(HASH(concat(issue_type, severity))) AS ticket_key, issue_type, severity
223 FROM tickets WHERE issue_type IS NOT NULL;
224
225 INSERT INTO TABLE Dim_Ticket VALUES (-1, 'No Ticket', 'N/A');
226

-- Load Dim_Product
228
229 INSERT OVERWRITE TABLE Dim_Product
230 SELECT DISTINCT ABS(HASH(product_type)) AS product_key, product_type
231 FROM `usage` WHERE product_type IS NOT NULL;
232
233 INSERT INTO TABLE Dim_Product VALUES (-1, 'No Product');
234
```

Add a name... Add a description...

0.95s project1_dw

```

228
229 INSERT OVERWRITE TABLE Dim_Product
230 SELECT DISTINCT ABS(HASH(product_type)) AS product_key, product_type
231 FROM `usage` WHERE product_type IS NOT NULL;
232
233 INSERT INTO TABLE Dim_Product VALUES (-1, 'No Product');
234
235 SELECT * from Dim_Product;
236
237
238 -- Dim_Date
239
240 INSERT OVERWRITE TABLE Dim_Date
241 SELECT DISTINCT

```

Query History

Saved Queries

Results (6)

COLUMNS (2) Q

dim_product.product_key

int

dim_product.product_type

varc

	dim_product.product_key	dim_product.product_type
1	2373904	Loan
2	372502797	Investment
3	757365607	Savings
4	1304940503	Credit Card
5	1601542650	Checking
6	-1	No Product

-- Load Dim_Date

```

237 |
238 INSERT OVERWRITE TABLE Dim_Date
239 SELECT DISTINCT
240     CAST(date_format(dt, 'yyyyMMdd') AS INT) AS date_key,
241     dt AS full_date,
242     day(dt) AS day,
243     month(dt) AS month,
244     date_format(dt, 'MMMM') AS month_name,
245     quarter(dt) AS quarter,
246     year(dt) AS year
247 FROM (
248     SELECT CAST(log_date AS DATE) AS dt FROM `usage` WHERE log_date IS NOT NULL
249     UNION
250     SELECT date_offered AS dt FROM offers WHERE date_offered IS NOT NULL
251     UNION
252     SELECT date_opened AS dt FROM customers WHERE date_opened IS NOT NULL
253 ) all_dates
254 WHERE dt IS NOT NULL;
255

```

```

-- =====
-- Loading Fact Table
-- =====

```

INSERT OVERWRITE TABLE Fact_Customer_Activity PARTITION(event_type, event_year)

-- 1. Usage Events

```

SELECT
    CAST(ABS(HASH(usage_log_id)) AS BIGINT) AS activity_key,
    ABS(HASH(customer_id)) AS cust_key,
    CAST(date_format(CAST(log_date AS DATE), 'yyyyMMdd') AS INT) AS date_key,
    -1 AS ticket_key, -1 AS offer_key, product_type,
    CAST(monthly_balance AS DECIMAL(15,2)) AS usage_balance,
    NULL AS resolution_time_hrs, NULL AS is_accepted,
    'Usage' AS event_type, year(CAST(log_date AS DATE)) AS event_year

```

FROM `usage`

UNION ALL

-- 2. Offer Events

```

SELECT
    CAST(ABS(HASH(concat(customer_id, offer_id))) AS BIGINT) AS activity_key,
    ABS(HASH(customer_id)) AS cust_key,
    CAST(date_format(date_offered, 'yyyyMMdd') AS INT) AS date_key,
    -1 AS ticket_key, ABS(HASH(offer_id)) AS offer_key, 'N/A' AS product_type,
    NULL AS usage_balance, NULL AS resolution_time_hrs, accepted AS is_accepted,
    'Offer' AS event_type, year(date_offered) AS event_year

```

FROM offers

UNION ALL

-- 3. Support Ticket Events

SELECT

```
CAST(ABS(HASH(ticket_id)) AS BIGINT) AS activity_key,  
ABS(HASH(customer_id)) AS cust_key, -1 AS date_key,  
ABS(HASH(ticket_id)) AS ticket_key, -1 AS offer_key, 'N/A' AS product_type,  
NULL AS usage_balance, CAST(resolution_time_hrs AS DECIMAL(10,2)), NULL AS is_accepted,  
'Support_Ticket' AS event_type, 1900 AS event_year
```

FROM tickets;


6. Data Extraction for BI Consumption

Finally, the analytical datasets were exported from Hive to the Hadoop Distributed File System (HDFS), and subsequently merged into flat .csv files on the local Linux file system for seamless ingestion into Power BI / Tableau.

Exporting Data from Hive to HDFS

SQL

```
317  
318 -- 1. Export Dim_Customer  
319 INSERT OVERWRITE DIRECTORY '/user/student/project1/gold/csv/Dim_Customer'  
320 ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
321 SELECT * FROM Dim_Customer;  
322  
323 -- 2. Export Dim Offer  
324 INSERT OVERWRITE DIRECTORY '/user/student/project1/gold/csv/Dim_Offer'  
325 ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
326 SELECT * FROM Dim_Offer;  
327  
328 -- 2. Export Dim_Product  
329 INSERT OVERWRITE DIRECTORY '/user/student/project1/gold/csv/Dim_Product'  
330 ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
331 SELECT * FROM dim_product;  
332  
333 -- 3. Export Dim Ticket  
334 INSERT OVERWRITE DIRECTORY '/user/student/project1/gold/csv/Dim_Ticket'  
335 ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
336 SELECT * FROM Dim_Ticket;  
337  
338 -- 4. Export Dim Date  
339 INSERT OVERWRITE DIRECTORY '/user/student/project1/gold/csv/Dim_Date'  
340 ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
341 SELECT * FROM Dim_Date;  
342  
343 -- 5. Export Fact Customer Activity  
344 INSERT OVERWRITE DIRECTORY '/user/student/project1/gold/csv/Fact_Customer_Activity'  
345 ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
346 SELECT * FROM Fact_Customer_Activity;  
347
```



DETAILED GOLD LAYER (DATA WAREHOUSE) IMPLEMENTATION FLOW

Spark SQL — Monthly Churn Rate

```
: churn_by_cohort = (df_customers
  .filter(F.col("date_opened").isNotNull() & F.col("is_churned").isNotNull())
  .withColumn("cohort_month", F.date_format("date_opened", "yyyy-MM"))
  .groupBy("cohort_month")
  .agg(
    F.count("*").alias("total_customers"),
    F.sum(F.col("is_churned").cast("int")).alias("churned_customers"))
  .withColumn("churn_rate_pct", F.round(F.col("churned_customers") / F.col("total_customers") * 100, 2))
  .orderBy("cohort_month"))

churn_by_cohort.show(30)
```

[Stage 365:=====> (180 + 6) / 200]

cohort_month	total_customers	churned_customers	churn_rate_pct
2016-03	5	1	20.0
2016-09	485	100	20.62
2017-03	3	1	33.33
2017-09	981	212	21.61
2018-03	1	1	100.0
2018-09	1024	196	19.14
2019-03	2	1	50.0
2019-09	1026	176	17.15
2020-02	2	0	0.0
2020-09	965	196	20.31
2021-02	6	1	16.67
2021-09	1006	208	20.68
2022-02	2	0	0.0
2022-09	987	203	20.57
2023-02	4	0	0.0
2023-09	1005	213	21.19
2024-01	3	0	0.0
2024-09	1045	201	19.23
2025-01	2	2	100.0
2025-09	1033	230	22.27
2026-01	2	0	0.0

Cohort Churn Rate by Signup Month

Target of this calculation: for each signup month (`date_opened`), what percentage of customers who opened their account that month are *currently* marked as churned (`is_churned` = 1). This is a **cohort snapshot**, not a time-series churn rate — it doesn't show churn *happening* in a given month, only the churn status *today* of everyone who joined in that month. True monthly churn (customers who left *during* that specific month) would need a churn-event date, which doesn't exist in this data.

What the output presents: one row per signup month, with `total_customers` in that cohort, how many of them are `churned_customers`, and the resulting `churn_rate_pct`.

Executive Customer Churn & Revenue Risk Dashboard: Technical & Business Report



Strategic Objective & Business Value

Designed and implemented an executive-level analytics solution tailored for engineering and product leadership to proactively monitor customer attrition, quantify revenue exposure, and evaluate **Customer Lifetime Value (CLV)**. The system transforms raw customer data into actionable business insights, connecting operational risk metrics with financial outcomes.

UI/UX Architecture & Layout System

Executive Navigation Panel: Implemented a fixed-width sidebar containing global interactive controls and brand identifiers, providing a consistent and cohesive enterprise-style interface.

Top KPI Tier: Positioned a horizontal KPI section at the top of the dashboard displaying four core indicators—**Total Customers (10,000)**, **Total Churned (2,037)**, **Churn Rate (20.4%)**, and **Average CLV (258,318)**—providing an immediate overview of overall customer and business health.

Structured 2×2 Analytical Grid: Organized the main analytical visualizations into a balanced four-quadrant layout to analyze demographic vulnerability, geographic patterns, customer retention milestones, and product adoption behavior.

Visual Structure & Aesthetic Design

Minimalist Presentation: Removed default chart grid lines to reduce visual clutter and improve data readability, keeping the focus on key trends and insights.

Clean UI Formatting: Applied structured container backgrounds and transparent header elements to achieve a clean, modern, and professional executive dashboard design.

Interactivity & Data Integration

Synchronized Global Filtering: Configured multi-variable, cross-source filters—including **Is Churned, Gender, Geography, and Products**—to work consistently across all connected dashboard views.

Real-Time Root-Cause Analysis: Enabled dynamic updates across multiple analytical views, allowing stakeholders to drill down into specific behavioral, demographic, and geographic segments to investigate potential drivers of customer churn.